

Launching v4.1.1 of the Artificial Analysis Intelligence Index

We have updated the Artificial Analysis Intelligence Index to v4.1.1 - this patch release upgrades our grader models, and brings the latest τ^3 -Banking version to Artificial Analysis

To keep the Artificial Analysis Intelligence Index the most useful synthesis metric for developers, we make regular updates to the included evaluations and our independent methodology. Today's update is a minor one to keep our existing evaluation set as reliable as possible.

Overall model rankings remain largely consistent, with a slight increase in scores due to improved grading robustness across the updated evaluations. Claude Opus 5 remains in the #1 position with an Index of 63.

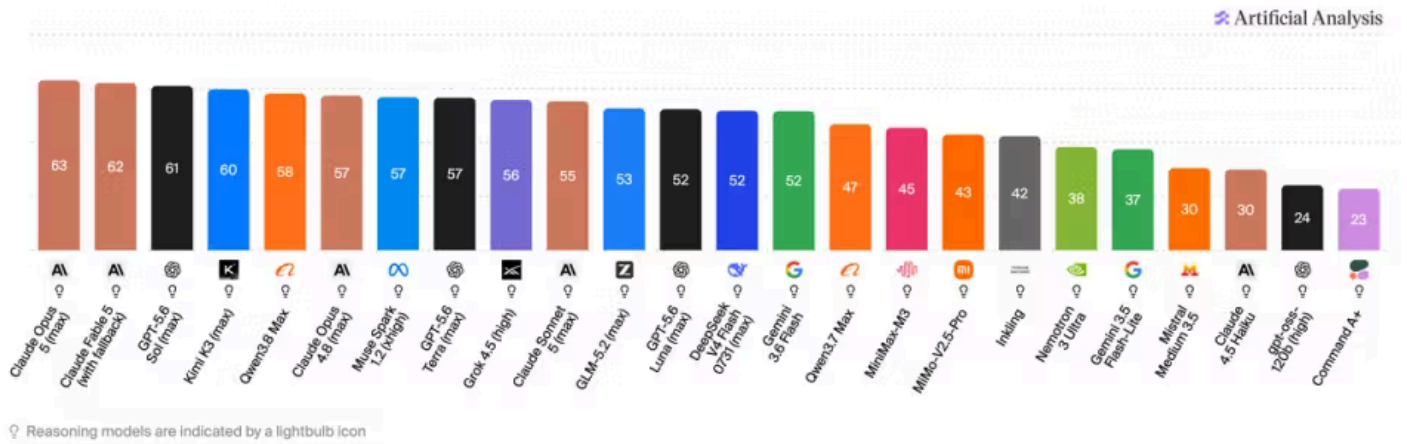
Key changes:

- ▶ τ^3 -Banking now runs v1.0.1 from Sierra, updating to the latest upstream task versions and improved grader pipeline that resolves correctness errors in trajectories that recover from unhappy paths
- ▶ HLE, AA-LCR and AA-Omniscience are now graded by GPT-5.6 Luna (medium), replacing GPT-4o, Qwen3 235B A22B 2507, and Gemini 3 Flash Preview respectively. These checks are now unified under a more capable modern model, selected for strong agreement with human judgment in our grader validation
- ▶ The effect on scores is small: most models move by less than a point on the Intelligence Index. The largest increase occurred for Muse Spark 1.2 (xhigh, +2.7 points), and the same models hold the top of the leaderboard

Published scores are now using v4.1.1, so all model results on Artificial Analysis now reflect these changes and use our latest consistent, independent methodology.

Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.1.1 incorporates 9 evaluations: GDPval-AA v2, r²-Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, AA-LCR



Newsletter


Get notified about new articles

What's your email?

Subscribe

We'll email you when we publish something new.


Read the latest



Agnes AI releases Agnes 2.5 Pro Beta

Agnes 2.5 Pro Beta


August 27, 2026



Intelligence at pocket scale: Benchmarking small models and mobile phones

Independent intelligence benchmarking of small language models on a set of evaluations chosen for mobile device use, launched alongside mobile phone inference benchmarking with Liquid AI. We evaluate the sam...


August 24, 2026



Announcing the Speech Agent Arena: Compare Speech agents in real world conversations

Announcing our new Speech Agent Arena, evaluating Speech to Speech models on real-world scenarios to analyze conversational preference and task success rate

August 24, 2026



Artificial Analysis

Get notified about new articles

Subscribe

Explore

[LLM Leaderboard](#)

[Image Arena](#)

[Video Arena](#)

[AI Agents](#)

[Evaluations](#)

Company

[About](#)

[Methodology](#)

[Services](#)

[Contact](#)

[Articles](#)

